



Departamento de Eletrónica, Telecomunicações e  
Informática

# **Machine Learning**

## **LECTURE 7: UNSUPERVISED LEARNING (K-MEANS CLUSTERING AND PCA)**

**Petia Georgieva**  
**(petia@ua.pt)**

# **Outline**

**1. K-means clustering**

**2. Data dimensionality reduction**

- data compression**
- data visualization**

**3. Principal Component Analysis (PCA)**

# SUPERVISED vs. UNSUPERVISED LEARNING

**Supervised Learning** - (given DATA + LABELS):

ML method is trained with labeled data to predict the labels of new examples (learning by labeled examples)

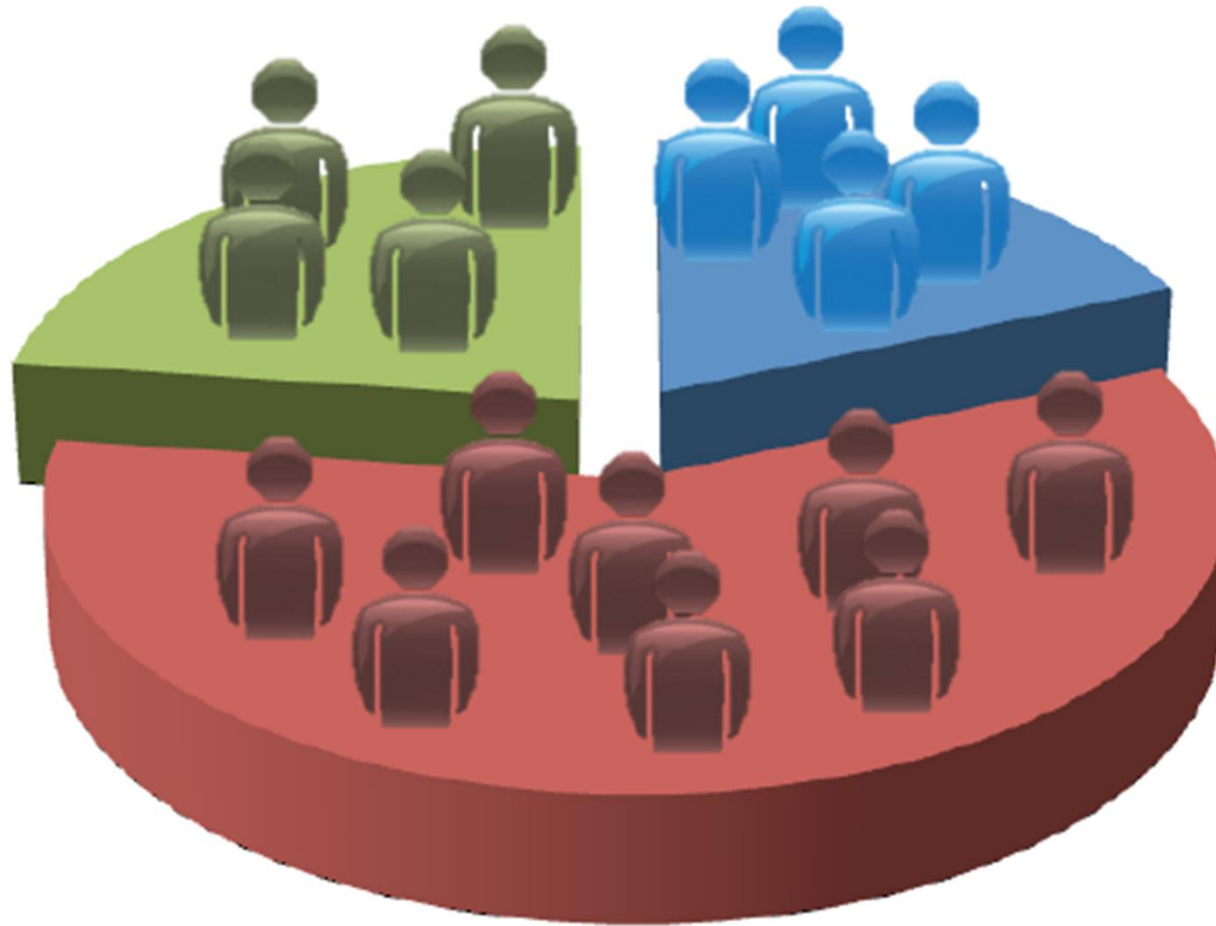
**Unsupervised Learning** - given UNLABELED DATA

ML method to discover the data internal structure

**Semi-Supervised Learning** – mixture of labeled and unlabeled data

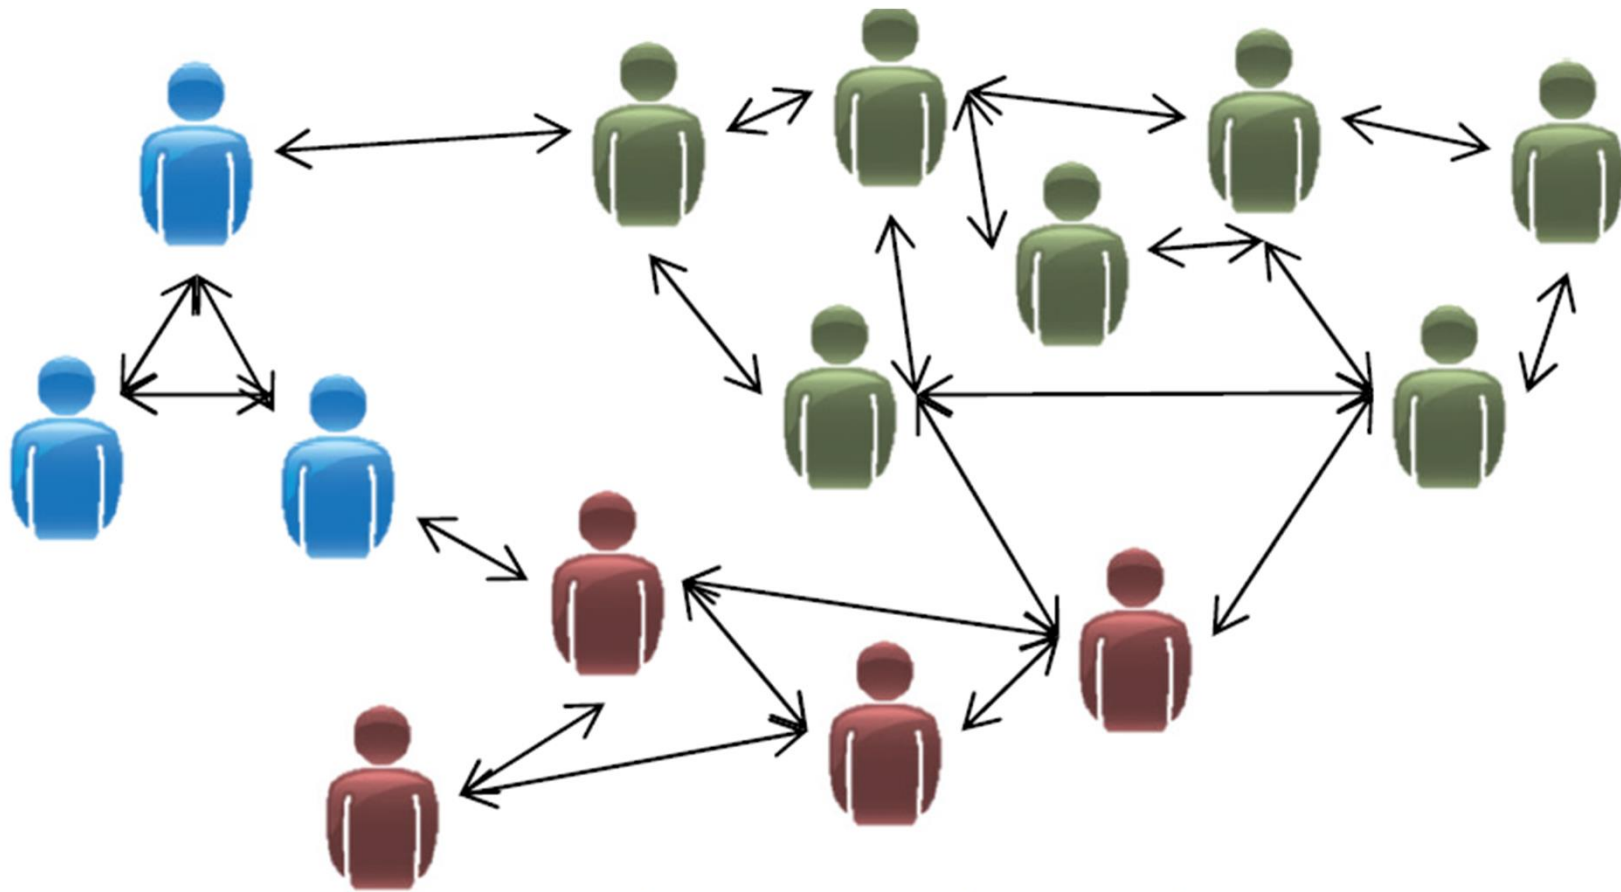
# Unsupervised learning -

**Market segmentation (data base of customers)**

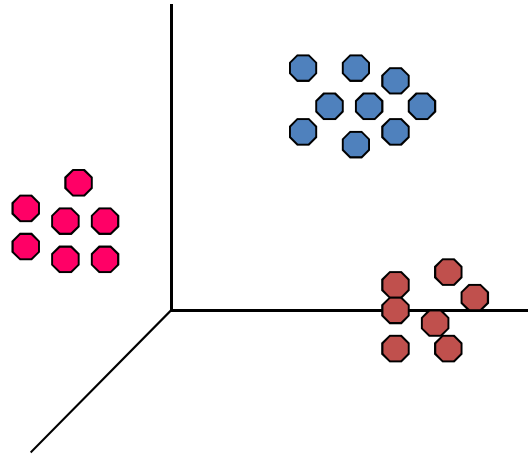


# Unsupervised learning

## Social network analysis



# Clustering intuition

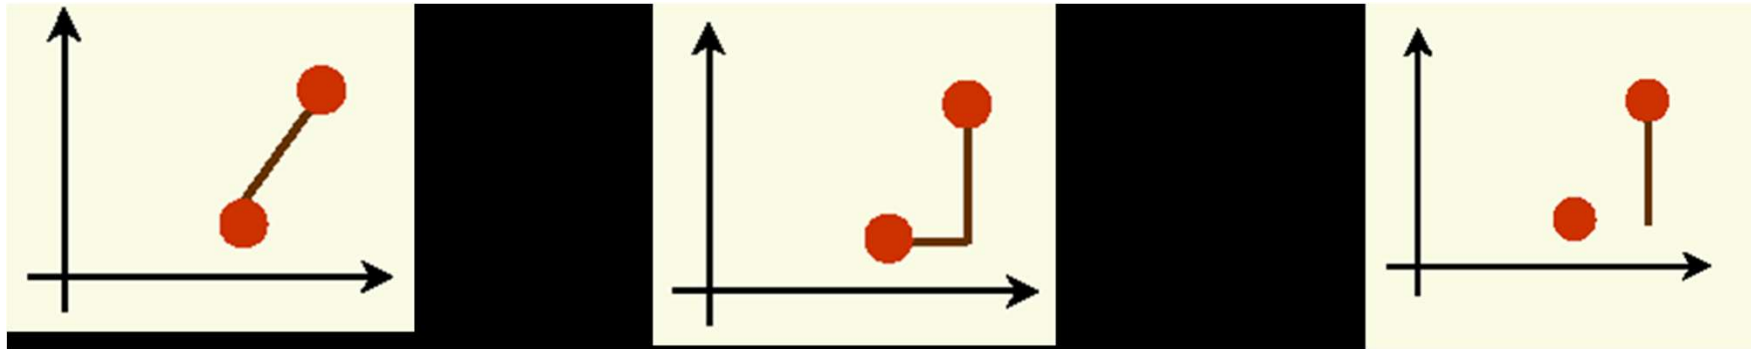


- Given a set of not labeled examples
- Find a relevant grouping of the examples into clusters such that:
  - Examples in the same cluster have **high similarity**
  - Examples from different clusters have **high dissimilarity**

## Similarity measures –

Euclidian distance; Chebyshev distance; Manhattan distance

# Distance (similarity) measures



**Euclidian Distance  
(L2 norm)**

$$d(p, q) = \sqrt{(x_p - x_q)^2 + (y_p - y_q)^2}$$

**Manhattan Distance  
(L1 norm)**

$$d(p, q) = |x_p - x_q| + |y_p - y_q|$$

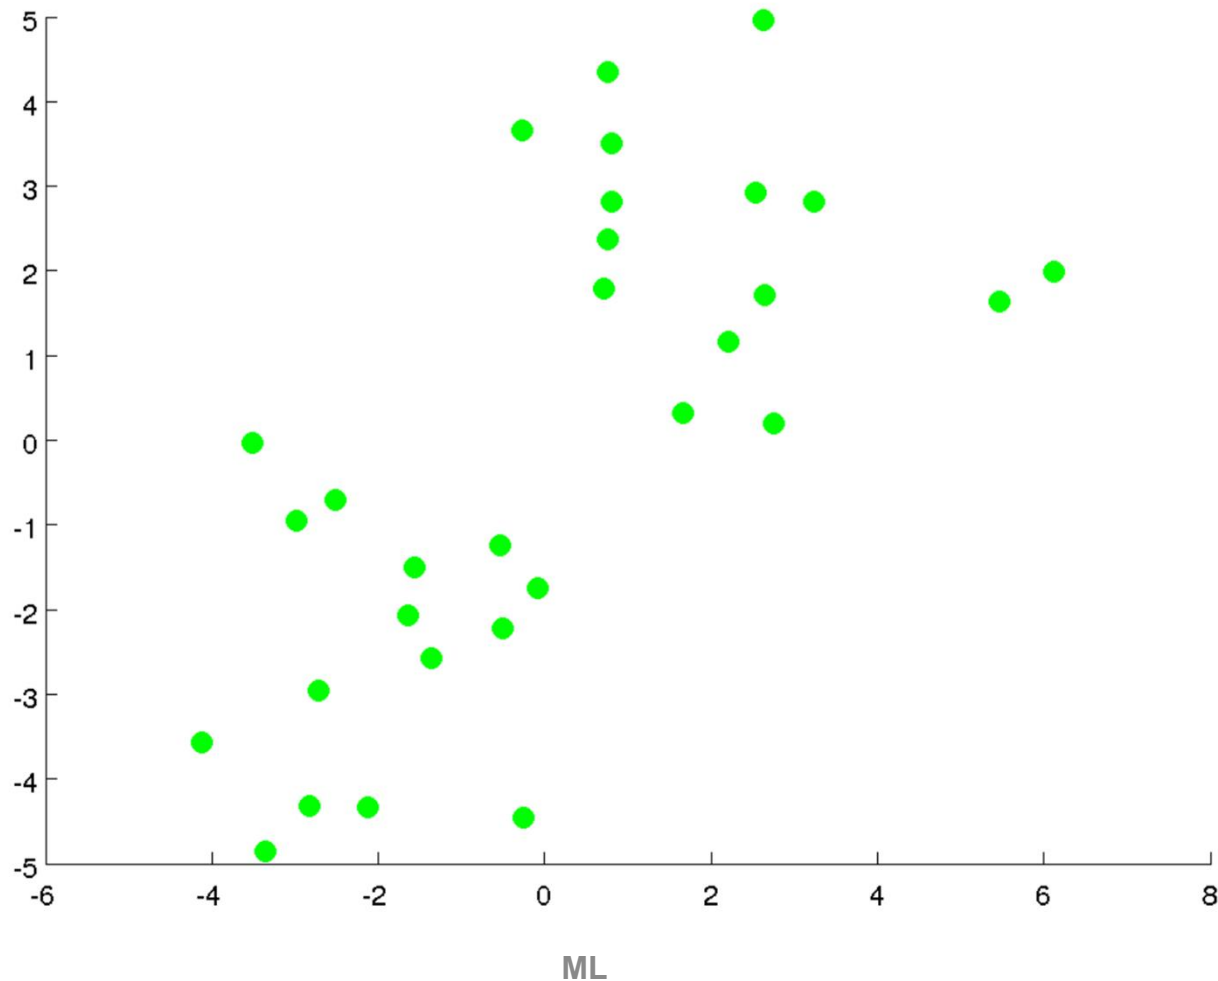
**Chebyshev distance**

$$d(p, q) = \max(|x_p - x_q|, |y_p - y_q|)$$

# K-means algorithm

**Given input:**

- **K (number of clusters)**
- **Training set (no labels)**

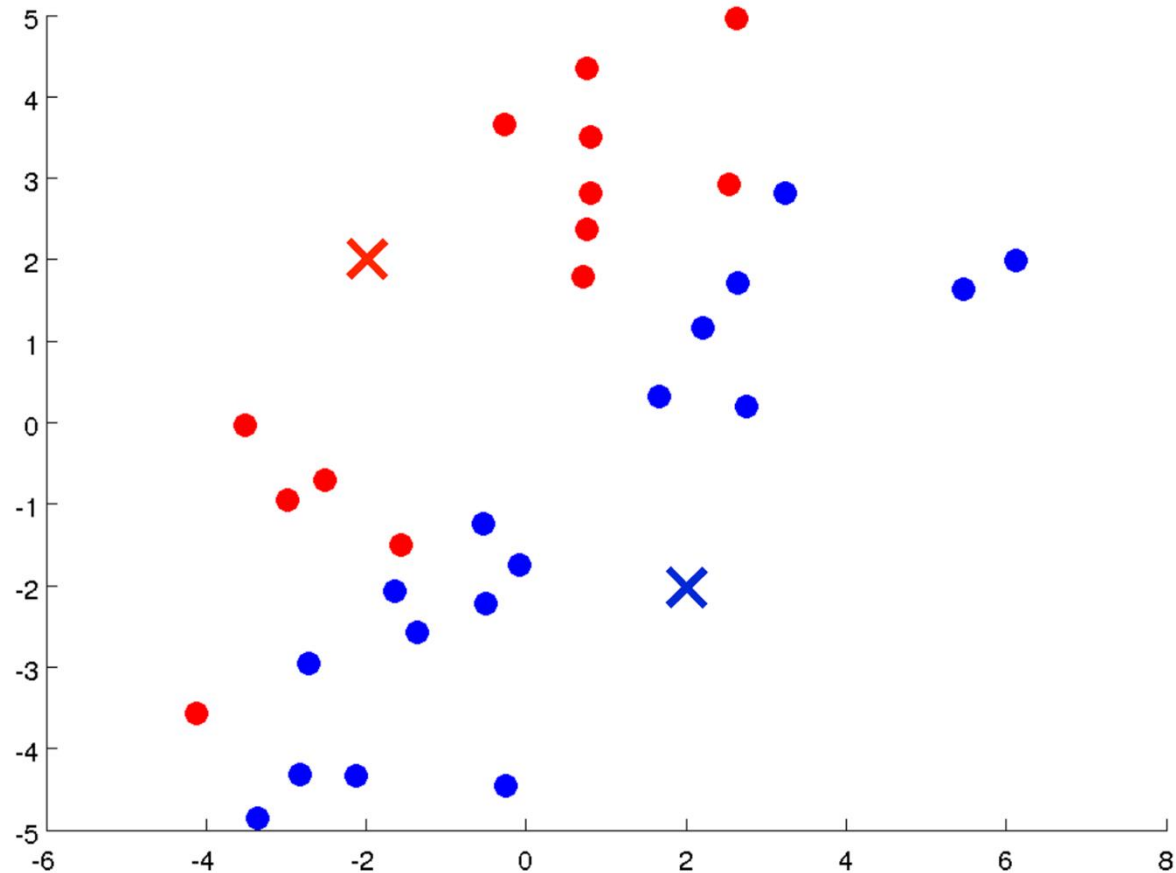




# K-means algorithm

**Randomly initialize K cluster centroids (e.g. K=2)**

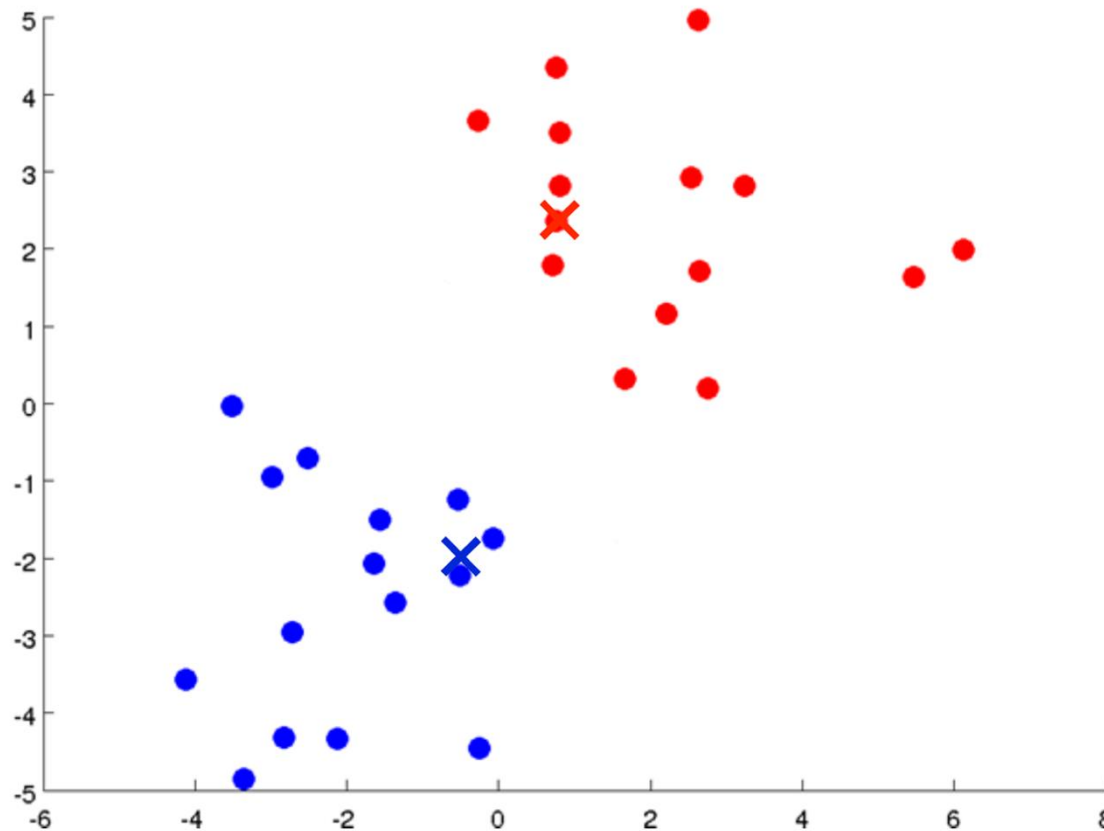
**Assign data points to their closest centroid (Euclidian distance)**



# K-means algorithm

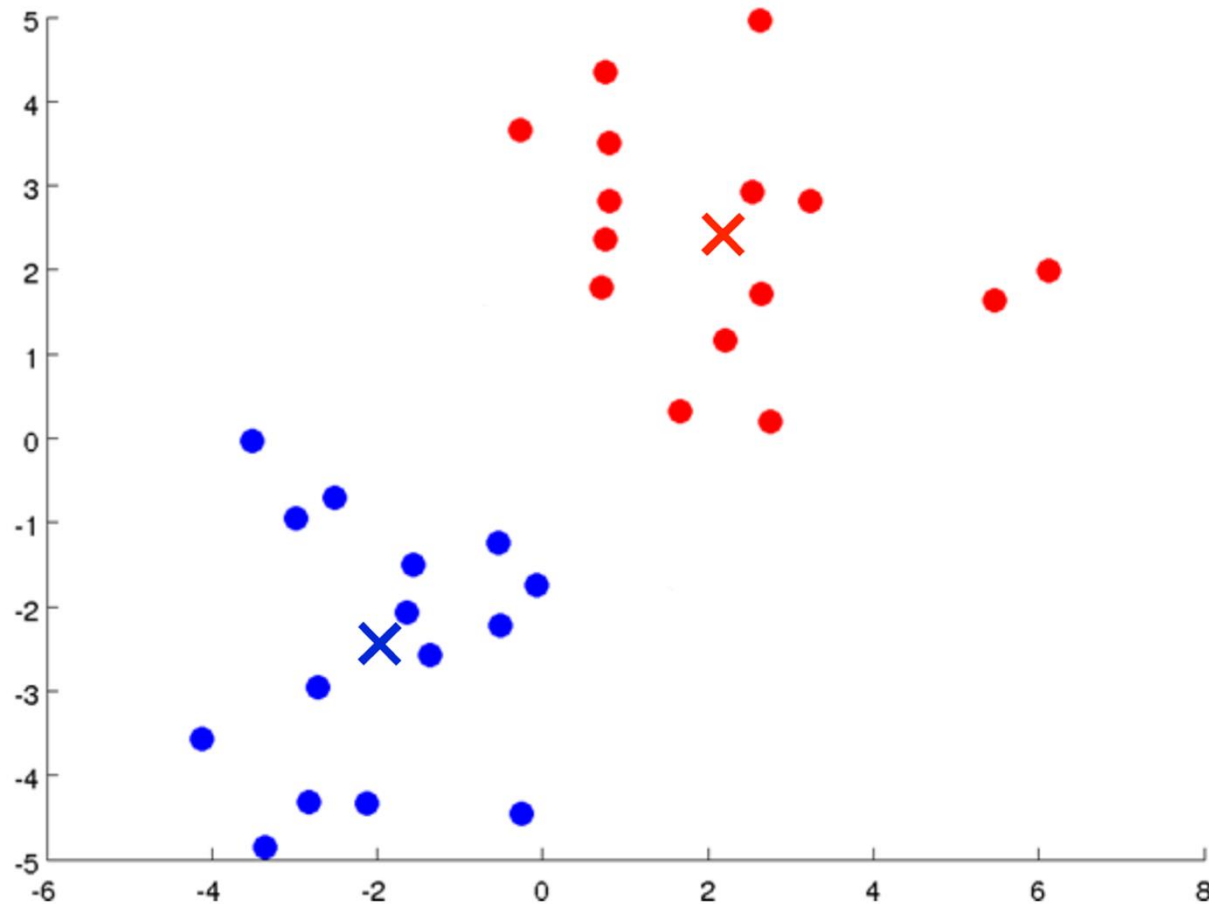
**Compute new centroids = mean of the points assigned to that cluster.**

**Assign data points to the new closest centroid.**



# K-means algorithm

Repeat until convergence



# K-means algorithms

Input:

- $K$  (number of clusters)
- Training set (no labels)

Randomly initialize  $K$  cluster centroids  $\mu_1, \mu_2, \dots, \mu_K \in \mathbb{R}^n$

Repeat {

for  $i = 1$  to  $m$

$c^{(i)} :=$  index (from 1 to  $K$ ) of cluster centroid  
closest to  $x^{(i)}$

for  $k = 1$  to  $K$

$\mu_k :=$  average (mean) of points assigned to cluster  $k$

}

*Cluster*

*assignment =>*  
*step*

*Move centroid =>*  
*step*

# K-means optimization objective (distortion = average distance)

$$J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K) = \frac{1}{m} \sum_{i=1}^m ||x^{(i)} - \mu_{c^{(i)}}||^2$$

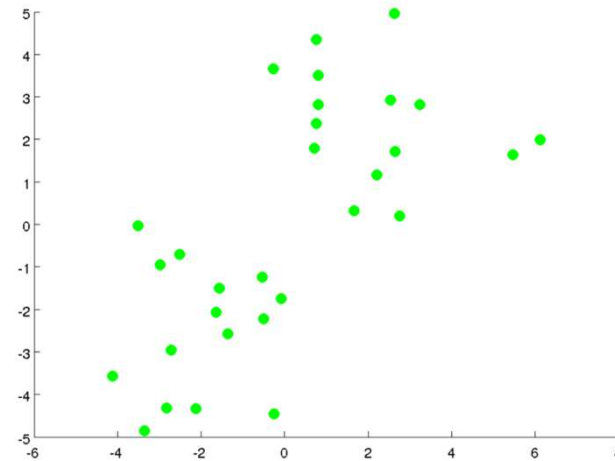
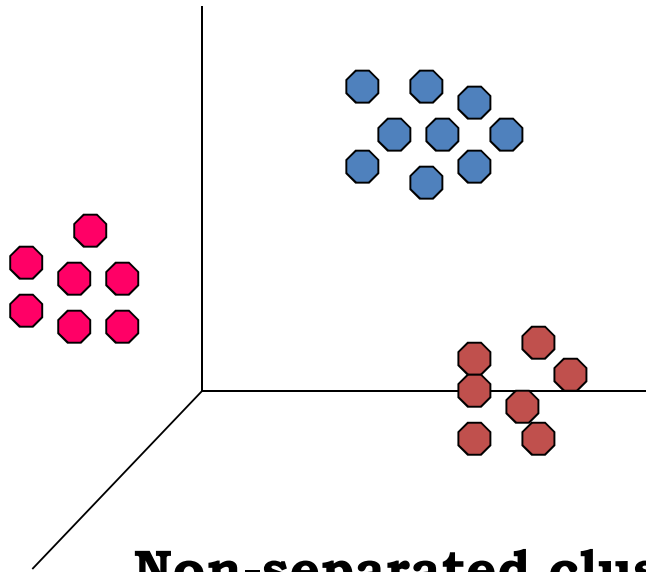
$$\min_{\substack{c^{(1)}, \dots, c^{(m)}, \\ \mu_1, \dots, \mu_K}} J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$$

Stop K-means learning (different criteria):

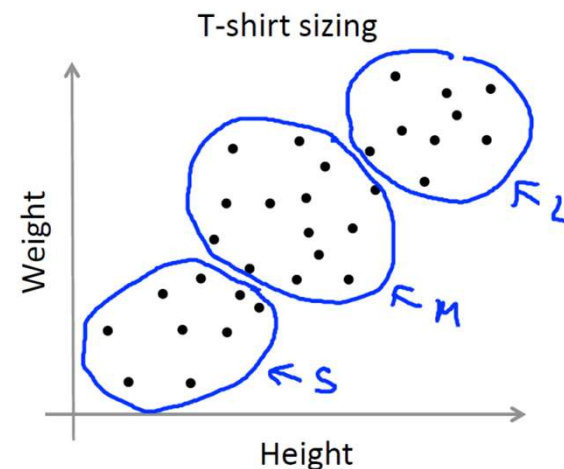
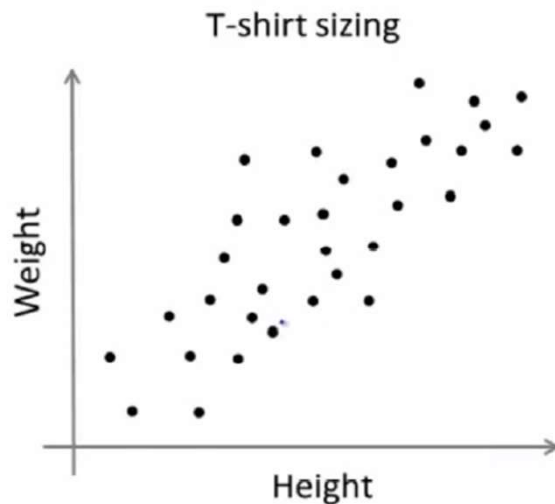
- Achieved Max number of iterations
- $J < \text{some threshold}$
- No improvement of  $J$  between  $x$  subsequent iterations

# K-means for non-separated clusters

## Separated clusters



**Non-separated clusters** (find problem specific criteria)

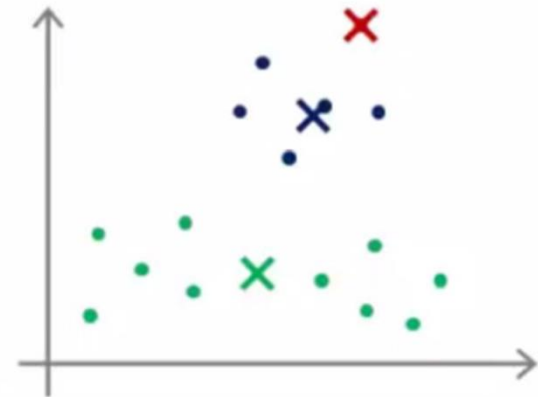
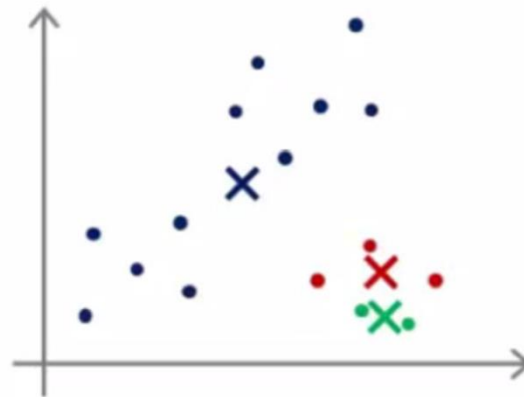
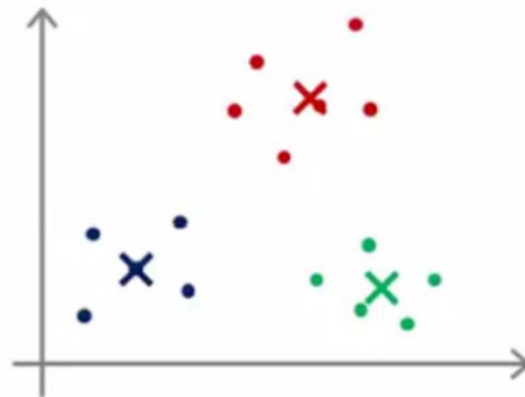
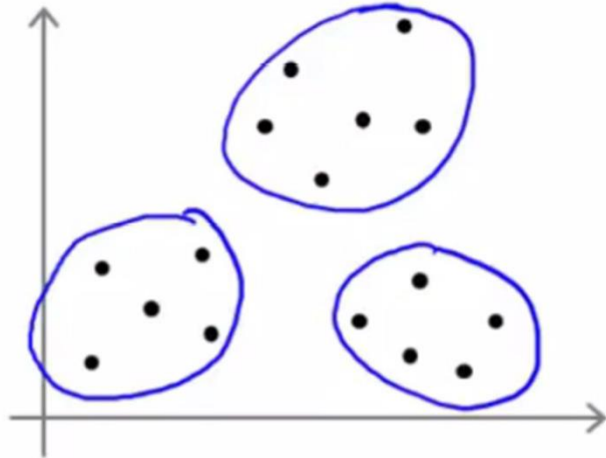


# Single (Random) Initialization

Choose # of clusters  $K < m$  (# of examples)

Initialize  $K$  cluster centroids = randomly picked  $K$  training examples

Local optima



# Repeat Random Initializations

For  $i = 1$  to 100 {

Randomly initialize K-means.

Run K-means. Get  $c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K$ .

Compute cost function (distortion)  
 $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$

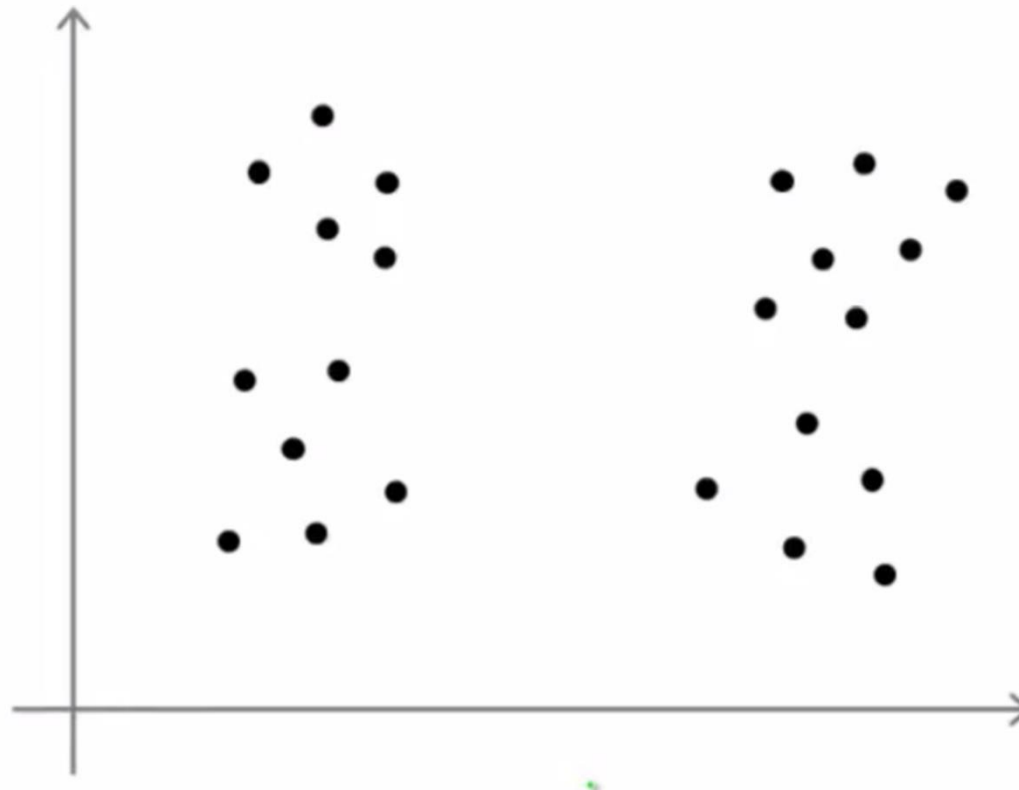
}

Pick clustering that gave lowest cost  $J(c^{(1)}, \dots, c^{(m)}, \mu_1, \dots, \mu_K)$

Repeat random inicializations works well for relatively small number of cluster e.g.  $K = (2, \dots, 10)$ .



# Choosing the number of clusters

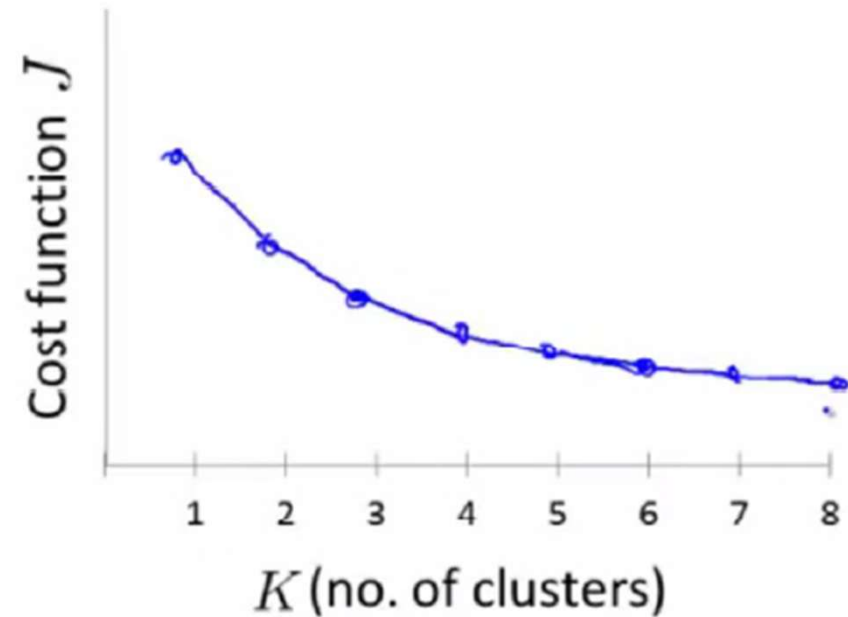
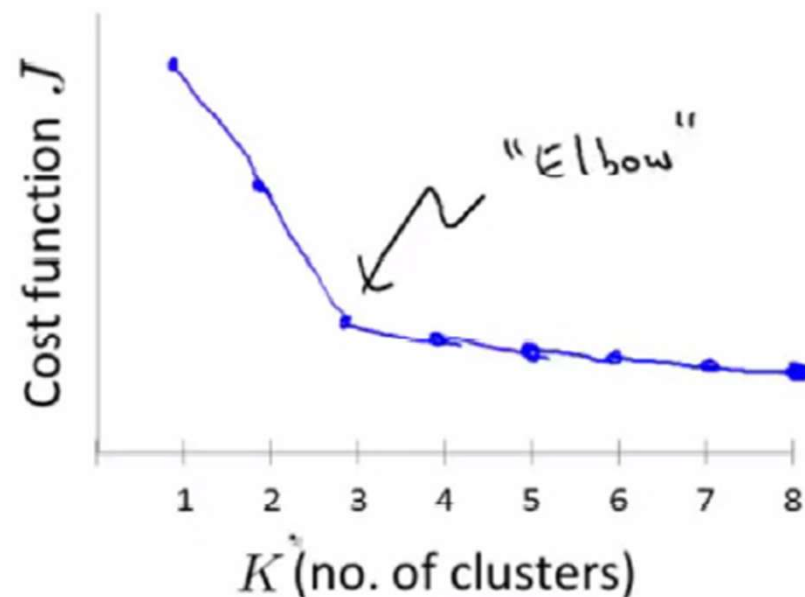


- Choose K by data visualisation (if possible)
- Ask domain experts (highly recommendable) , e.g. anomaly detection (experts should know how many types of anomalies are expected)
- Choose K automatically (e.g. Elbow method)

# Choosing the number of clusters (Elbow method)

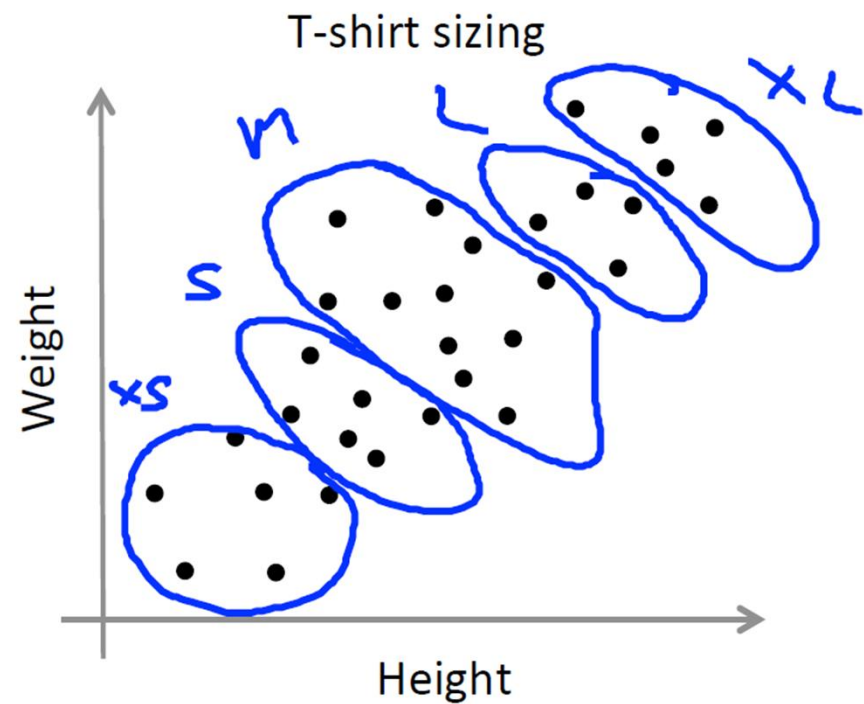
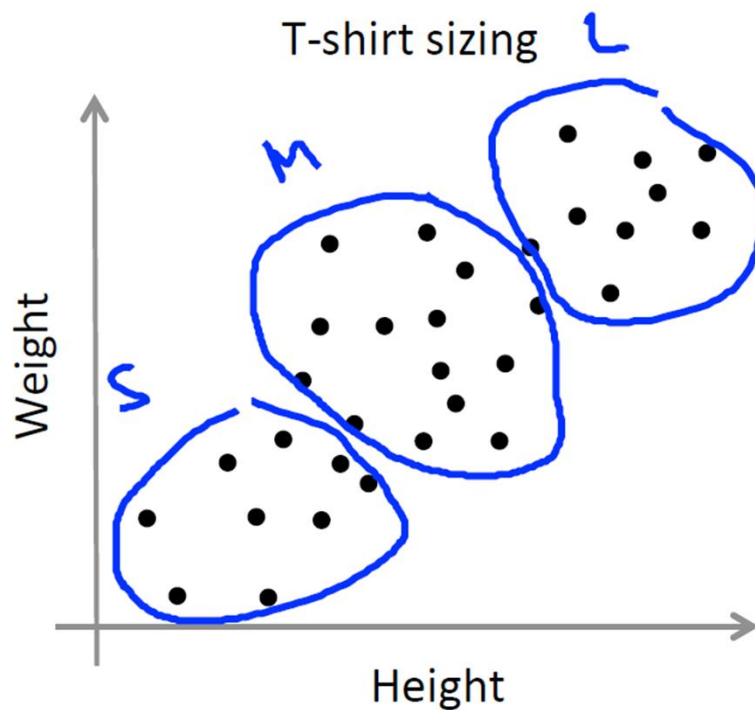
Appropriate for Elbow method

Not Appropriate for Elbow method



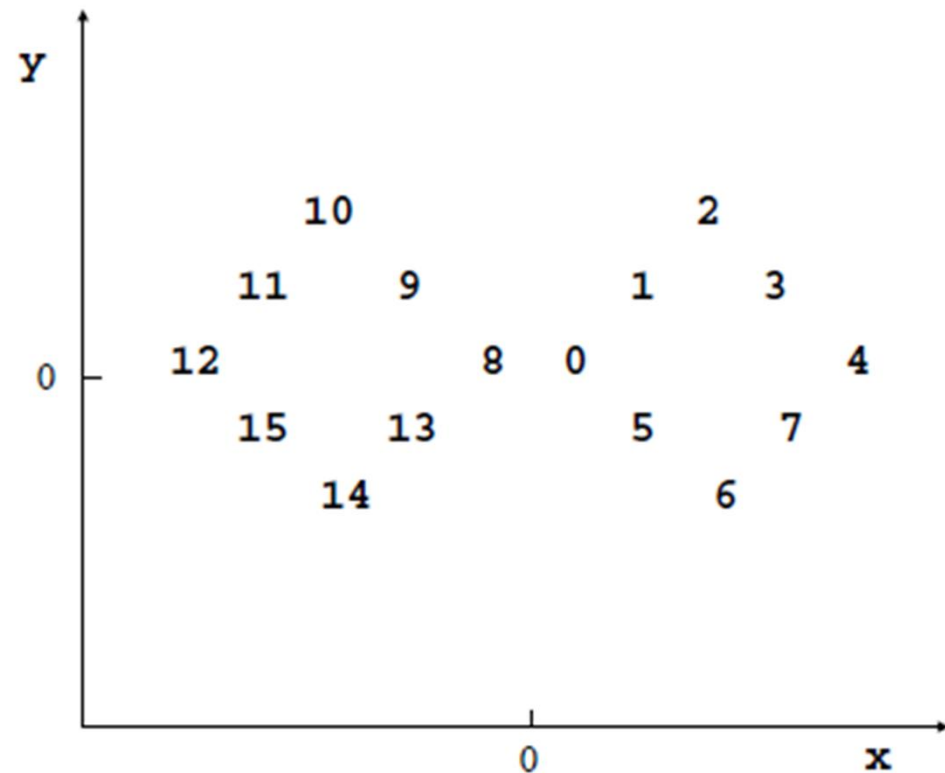
# Choosing the number of clusters

Run K-means to get clusters according to a particular purpose.  
(S, M, L type of t-shirts) or (XS, S, M, L, XL, of t-shirts) )



# K-MEANS CLUSTERING – Example

Id	x	y
0:	1.0	0.0
1:	3.0	2.0
2:	5.0	4.0
3:	7.0	2.0
4:	9.0	0.0
5:	3.0	-2.0
6:	5.0	-4.0
7:	7.0	-2.0
8:	-1.0	0.0
9:	-3.0	2.0
10:	-5.0	4.0
11:	-7.0	2.0
12:	-9.0	0.0
13:	-3.0	-2.0
14:	-5.0	-4.0
15:	-7.0	-2.0



- find the best 2 clusters

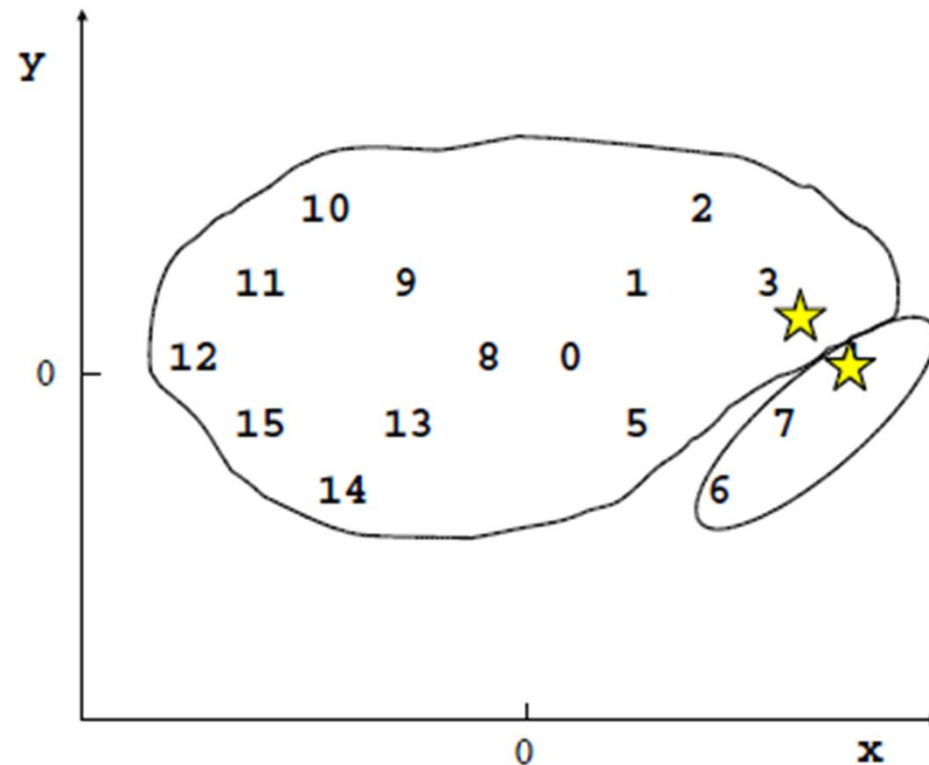
# K-MEANS CLUSTERING – Example

Seed: (9 0) (8 1)

Clustering: (4 6 7) (0 1 2 3 5 8 9 10 11 12 13 14 15)

Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

Average Distance: 4.35887



# K-MEANS CLUSTERING – Example

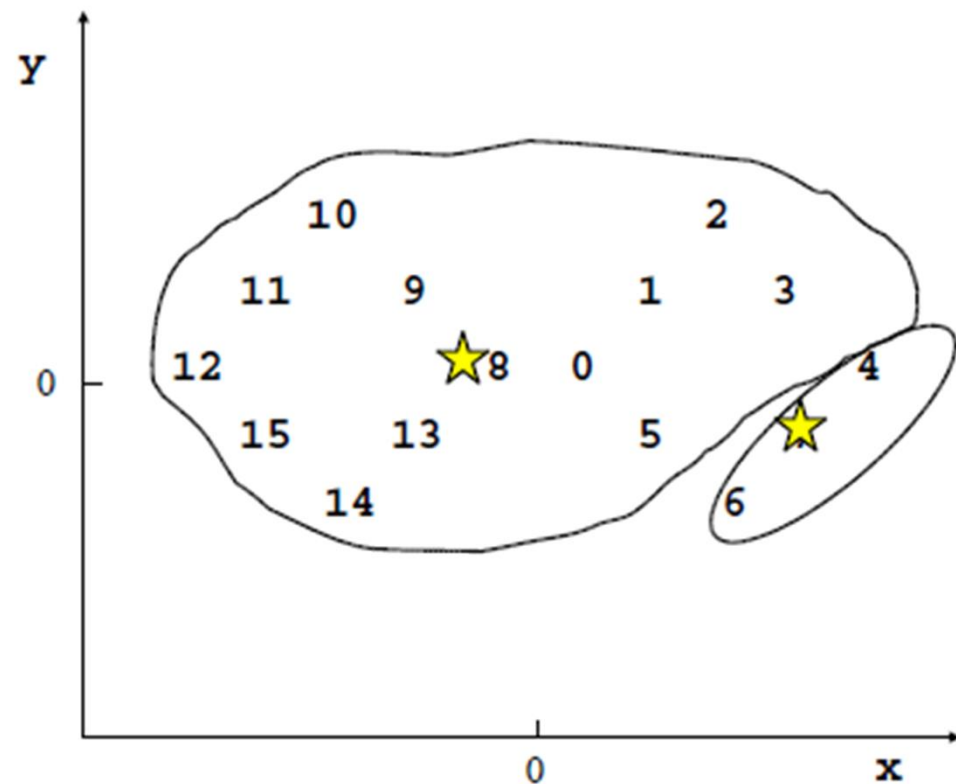
Seed: (9 0) (8 1)

Clustering: ( 4 6 7 ) ( 0 1 2 3 5 8 9 10 11 12 13 14 15 )

Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

Average Distance: 4.35887

Clustering: ( 2 3 4 5 6 7 ) ( 0 1 8 9 10 11 12 13 14 15 )



# K-MEANS CLUSTERING – Example

Seed: (9 0) (8 1)

Clustering: ( 4 6 7 ) ( 0 1 2 3 5 8 9 10 11 12 13 14 15 )

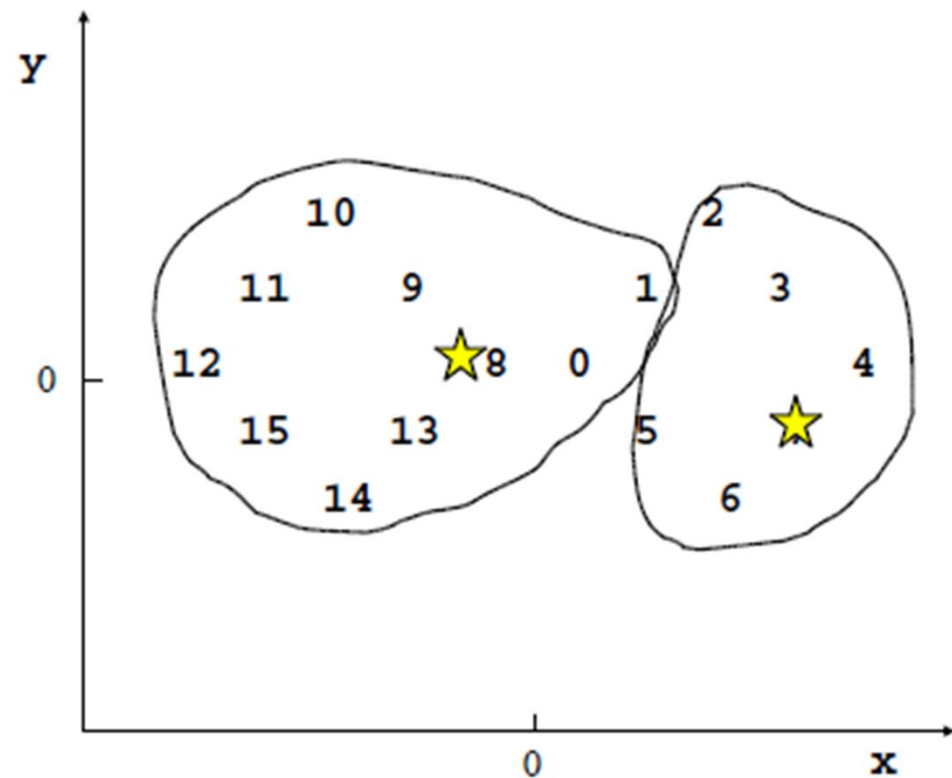
Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

Average Distance: 4.35887

Clustering: ( 2 3 4 5 6 7 ) ( 0 1 8 9 10 11 12 13 14 15 )

Cluster Centers: (6.0 -0.33334) (-3.6 0.2)

Average Distance: 3.6928



# K-MEANS CLUSTERING – Example

Seed: (9 0) (8 1)

Clustering: ( 4 6 7 ) ( 0 1 2 3 5 8 9 10 11 12 13 14 15 )

Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

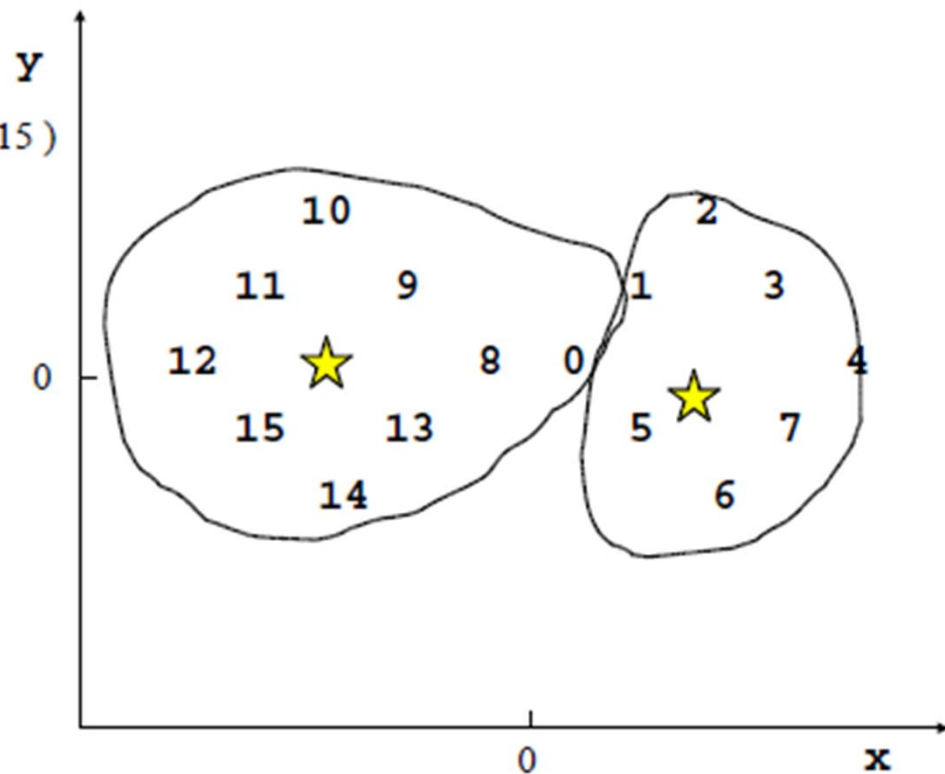
Average Distance: 4.35887

Clustering: ( 2 3 4 5 6 7 ) ( 0 1 8 9 10 11 12 13 14 15 )

Cluster Centers: (6.0 -0.33334) (-3.6 0.2)

Average Distance: 3.6928

Clustering: ( 1 2 3 4 5 6 7 ) ( 0 8 9 10 11 12 13 14 15 )





# K-MEANS CLUSTERING – Example

Seed: (9 0) (8 1)

Clustering: (4 6 7) (0 1 2 3 5 8 9 10 11 12 13 14 15)

Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

Average Distance: 4.35887

Clustering: (2 3 4 5 6 7) (0 1 8 9 10 11 12 13 14 15)

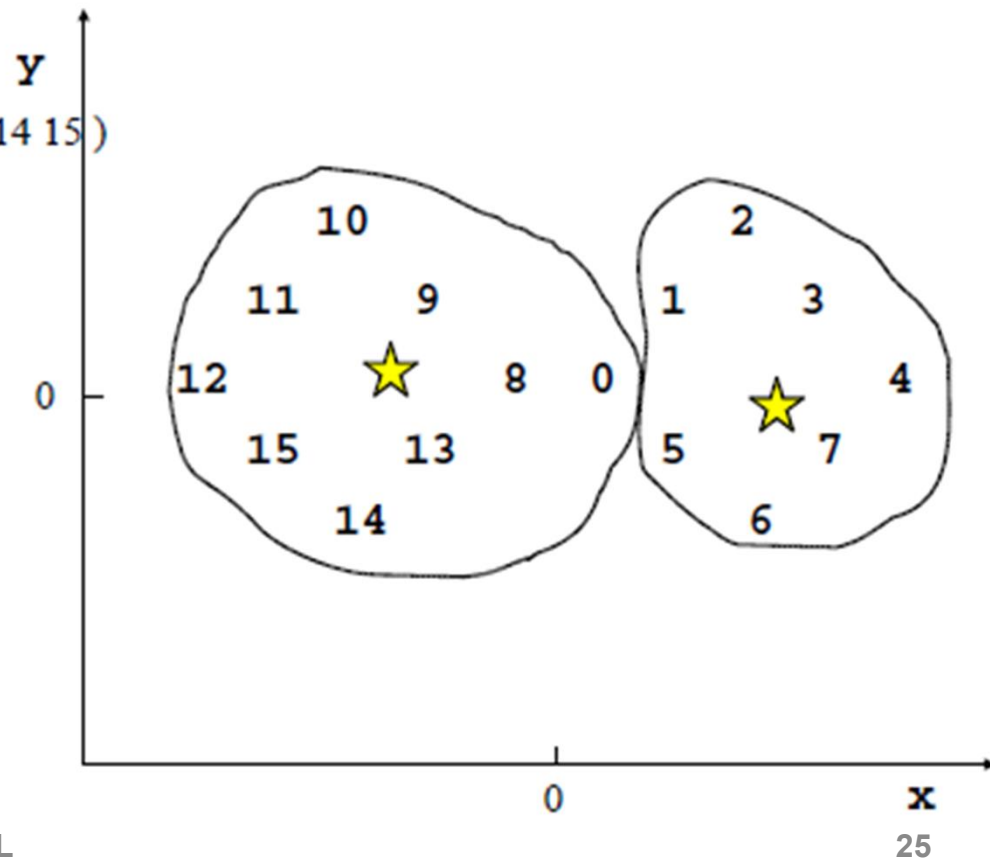
Cluster Centers: (6.0 -0.33334) (-3.6 0.2)

Average Distance: 3.6928

Clustering: (1 2 3 4 5 6 7) (0 8 9 10 11 12 13 14 15)

Cluster Centers: (5.57143 0.0) (-4.33334 0.0)

Average Distance: 3.49115



# K-MEANS CLUSTERING – Example

Seed: (9 0) (8 1)

Clustering: (4 6 7) (0 1 2 3 5 8 9 10 11 12 13 14 15)

Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

Average Distance: 4.35887

Clustering: (2 3 4 5 6 7) (0 1 8 9 10 11 12 13 14 15)

Cluster Centers: (6.0 -0.33334) (-3.6 0.2)

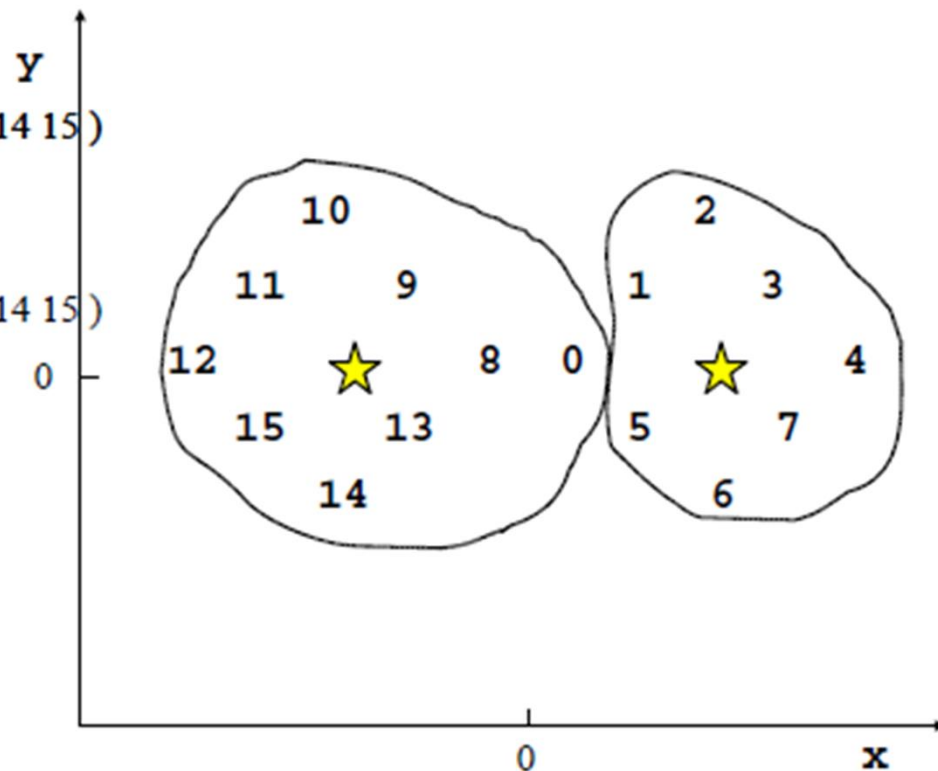
Average Distance: 3.6928

Clustering: (1 2 3 4 5 6 7) (0 8 9 10 11 12 13 14 15)

Cluster Centers: (5.57143 0.0) (-4.33334 0.0)

Average Distance: 3.49115

Clustering: (0 1 2 3 4 5 6 7) (8 9 10 11 12 13 14 15)



# K-MEANS CLUSTERING – Example

Seed: (9 0) (8 1)

Clustering: (4 6 7) (0 1 2 3 5 8 9 10 11 12 13 14 15)

Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

Average Distance: 4.35887

Clustering: (2 3 4 5 6 7) (0 1 8 9 10 11 12 13 14 15)

Cluster Centers: (6.0 -0.33334) (-3.6 0.2)

Average Distance: 3.6928

Clustering: (1 2 3 4 5 6 7) (0 8 9 10 11 12 13 14 15)

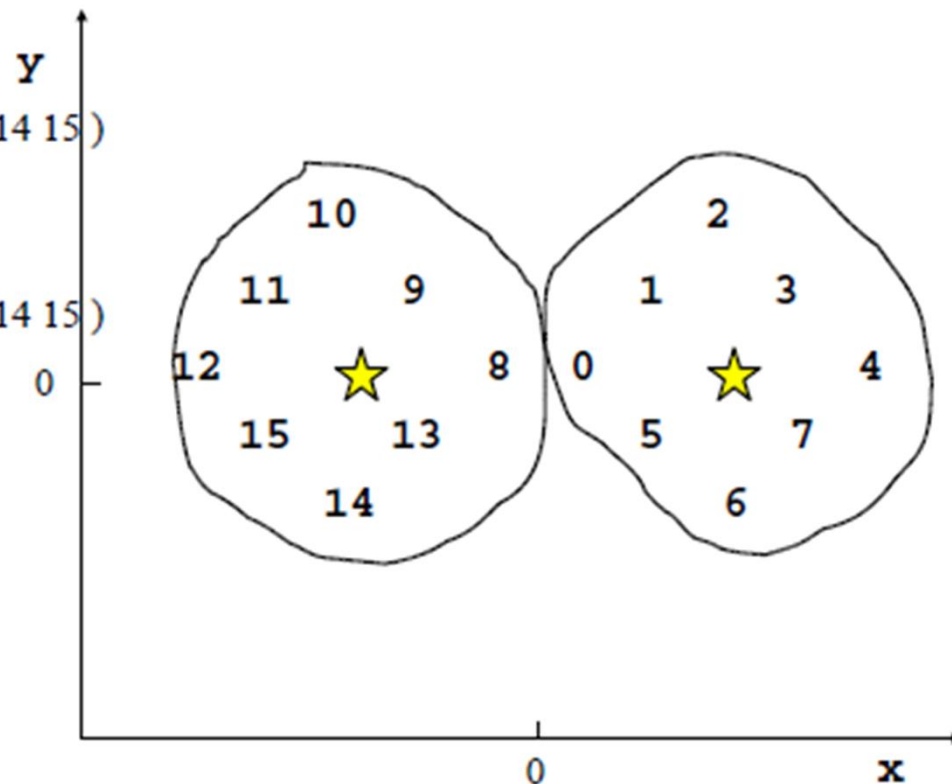
Cluster Centers: (5.57143 0.0) (-4.33334 0.0)

Average Distance: 3.49115

Clustering: (0 1 2 3 4 5 6 7) (8 9 10 11 12 13 14 15)

Cluster Centers: (5.0 0.0) (-5.0 0.0)

Average Distance: 3.41421



# K-MEANS CLUSTERING – Example

Seed: (9 0) (8 1)

Clustering: (4 6 7) (0 1 2 3 5 8 9 10 11 12 13 14 15)

Cluster Centers: (7.0 -2.0) (-1.61538 0.46153)

Average Distance: 4.35887

Clustering: (2 3 4 5 6 7) (0 1 8 9 10 11 12 13 14 15)

Cluster Centers: (6.0 -0.33334) (-3.6 0.2)

Average Distance: 3.6928

Clustering: (1 2 3 4 5 6 7) (0 8 9 10 11 12 13 14 15)

Cluster Centers: (5.57143 0.0) (-4.33334 0.0)

Average Distance: 3.49115

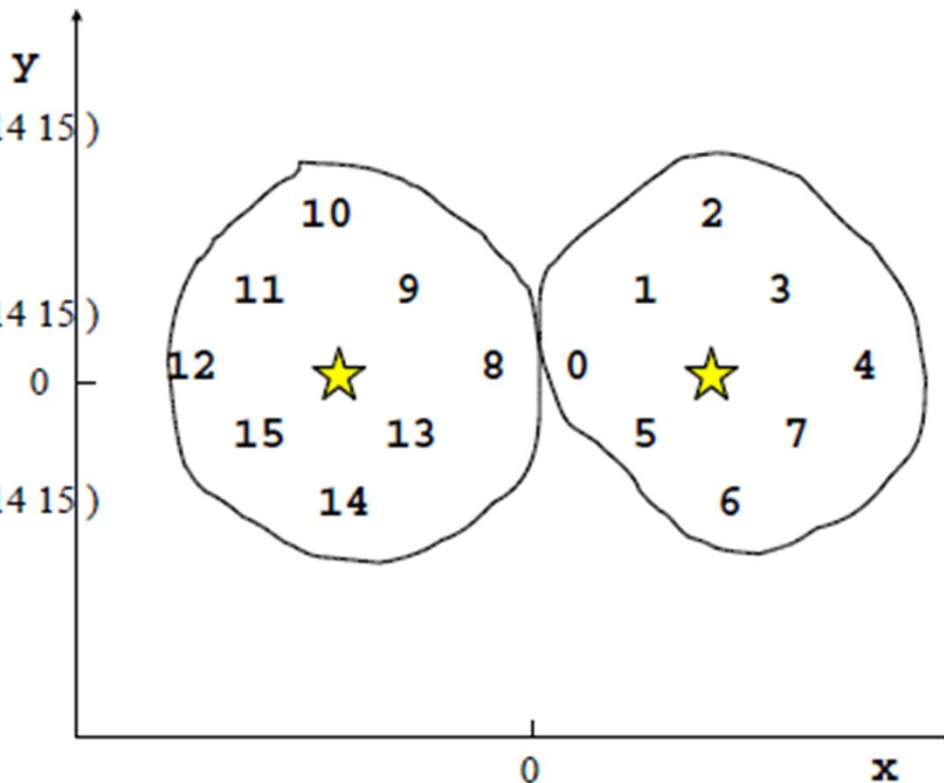
Clustering: (0 1 2 3 4 5 6 7) (8 9 10 11 12 13 14 15)

Cluster Centers: (5.0 0.0) (-5.0 0.0)

Average Distance: 3.41421

Clustering: (0 1 2 3 4 5 6 7) (8 9 10 11 12 13 14 15)

No improvement.



# HIERARCHICAL CLUSTERING

- Two approaches

- **Top-down:**

- start with 1 cluster (all examples belong to one cluster)
    - successively split into new sub clusters (based on some threshold max distance between the centroid and the examples)

- **Bottom-up:**

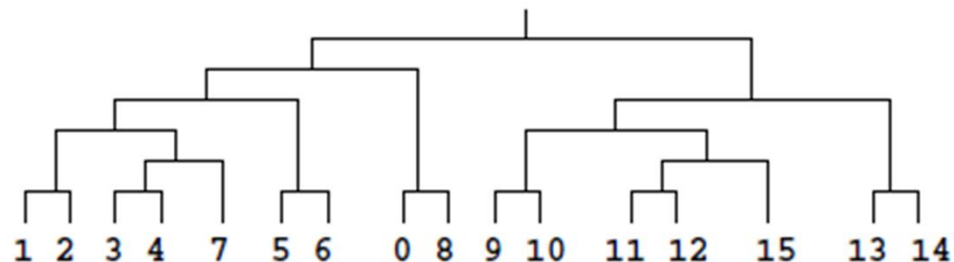
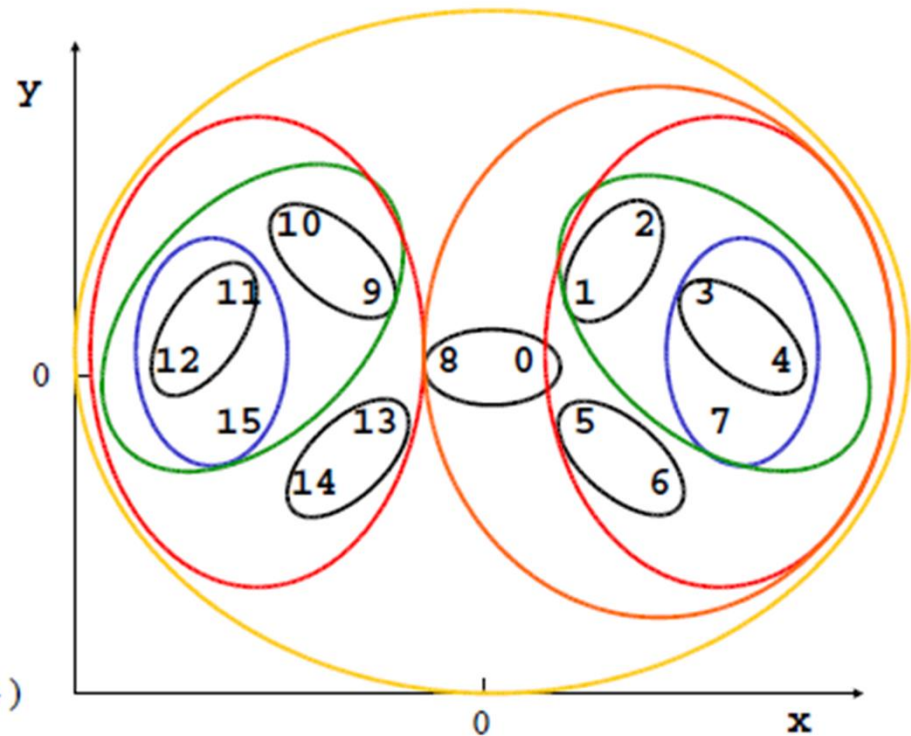
- start with as many clusters as examples
    - successively merge existing clusters  
(based on some similarity measure between clusters)



# SIMILARITY BETWEEN CLUSTERS

Bottom-up clustering (average-link):

min distance = 2.00000 (8)(0)  
 min distance = 2.82843 (2)(1)  
 min distance = 2.82843 (4)(3)  
 min distance = 2.82843 (6)(5)  
 min distance = 2.82843 (10)(9)  
 min distance = 2.82843 (12)(11)  
 min distance = 2.82843 (14)(13)  
 min distance = 3.16228 (7)(3 4)  
 min distance = 3.16228 (15)(11 12)  
 min distance = 4.73756 (3 4 7)(1 2)  
 min distance = 4.73756 (11 12 15)(9 10)  
 min distance = 4.74131 (1 2 3 4 7)(5 6)  
 min distance = 4.74131 (9 10 11 12 15)(13 14)  
 min distance = 5.57143 (0 8)(5 6 1 2 3 4 7)  
 min distance = 9.90476 (13 14 9 10 11 12 15)(5 6 1 2 3 4 7 0 8)



ML

# K-MEANS -summary

- The most popular clustering method.
- Need to know K.
- May converge to a Local Minimum .
- High number of computations.

# Expectation Maximization (EM)

**Expectation Maximization (EM)** is a generalization of K-Means.

Probabilistic approach, that assumes data from each cluster has a certain distribution (for example Gaussian Distribution).

Search for the Gaussian Distribution parameters (mean and standard deviation) that maximise the likelihood of data.

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



# Hard vs. Fuzzy CLUSTERING

**Hard Clustering**: Each point belongs to a single cluster.

**Fuzzy clustering**: Each example  $x_i$  belongs to any cluster  $c_j$  up to some degree depending on the values of membership function  $M_{cj}(x_i)$  .

- If  $M_{cj}(x_i)$  close to 1, the membership of  $x_i$  to cluster  $c_j$  is high
- If  $M_{cj}(x_i)$  close to 0, the membership of  $x_i$  to cluster  $c_j$  is low

# K-means for dimensionality reduction

dimensionality reduction is useful for:

## **Data compression (from 10000-1000 D to) 100D**

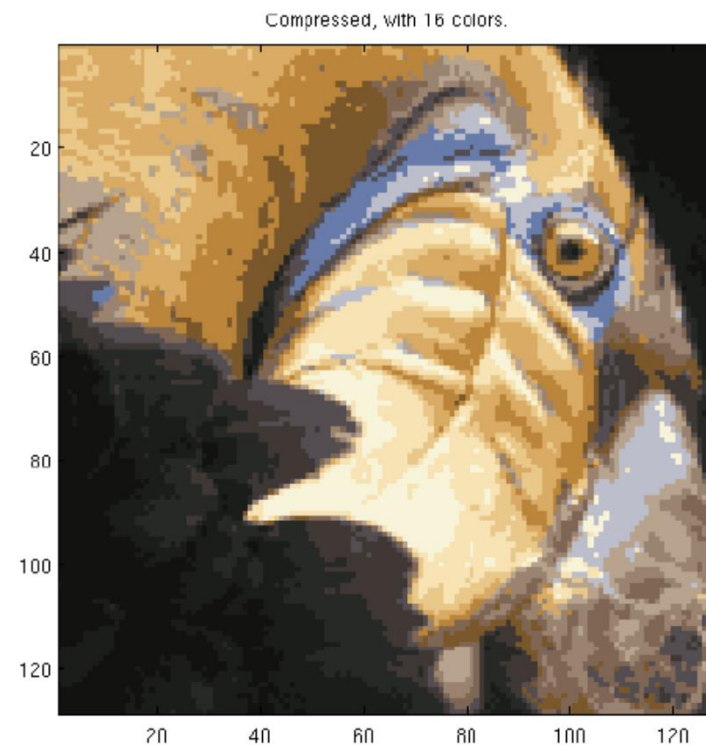
Reduce memory/ disk needed to store data  
Speed up learning algorithm

## **Data visualization (from 100-50D to 2-3D)**

**Remark:** kernel methods do the opposite, from the original low dimensional features (nonlinear models) go to a higher dimensional feature space (linear models)

# Image compression with K-means

RGB image:  $3 \times 8$  bits/pixel    Compressed image: 16 colors(clusters)  $\Rightarrow$  4 bits



# SEMI-SUPERVISED LEARNING

## Google Street View and clustering

**How to extrapolate concepts and label new objects in millions of unlabeled pictures. Feed a massive database of pictures to a machine, and use clustering to understand what is composed of, what are the elements in the pictures.**



# SEMI-SUPERVISED LEARNING

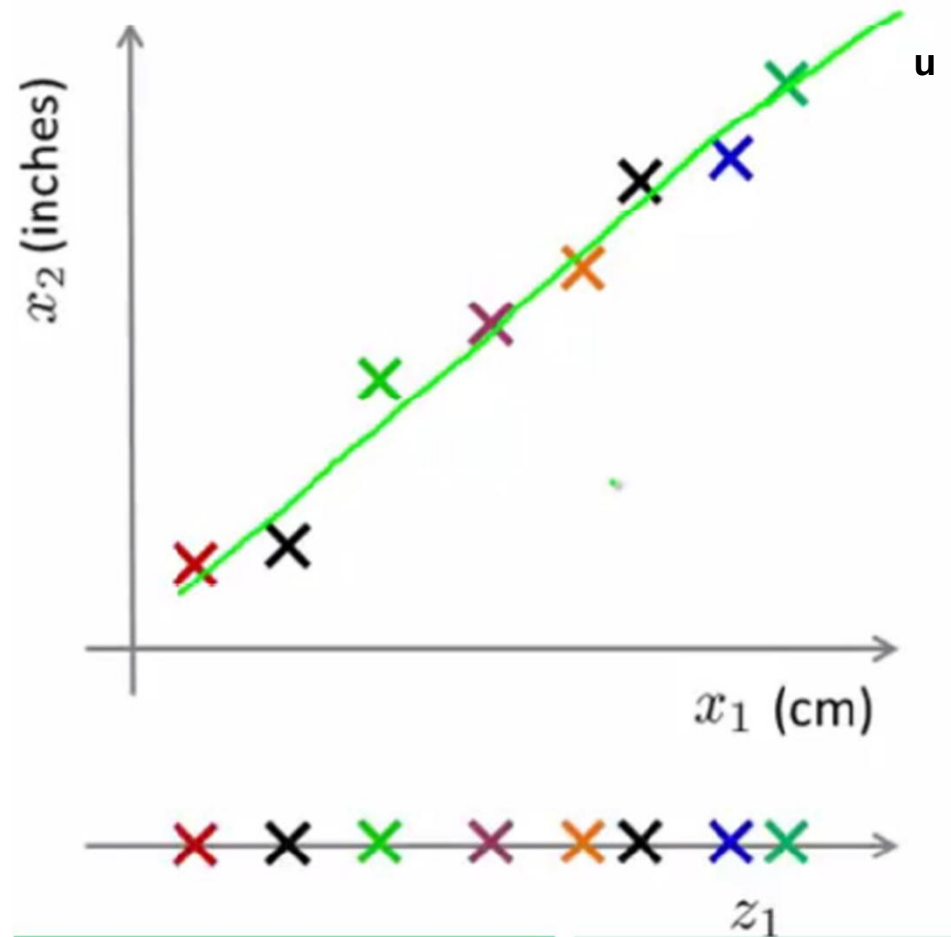
DARPA self-driving car <http://cs.stanford.edu/group/roadrunner/>



# DATA COMPRESSION

## Example: reduce data from 2D to 1D

If 2D data is located along a line, the second dimension is redundant

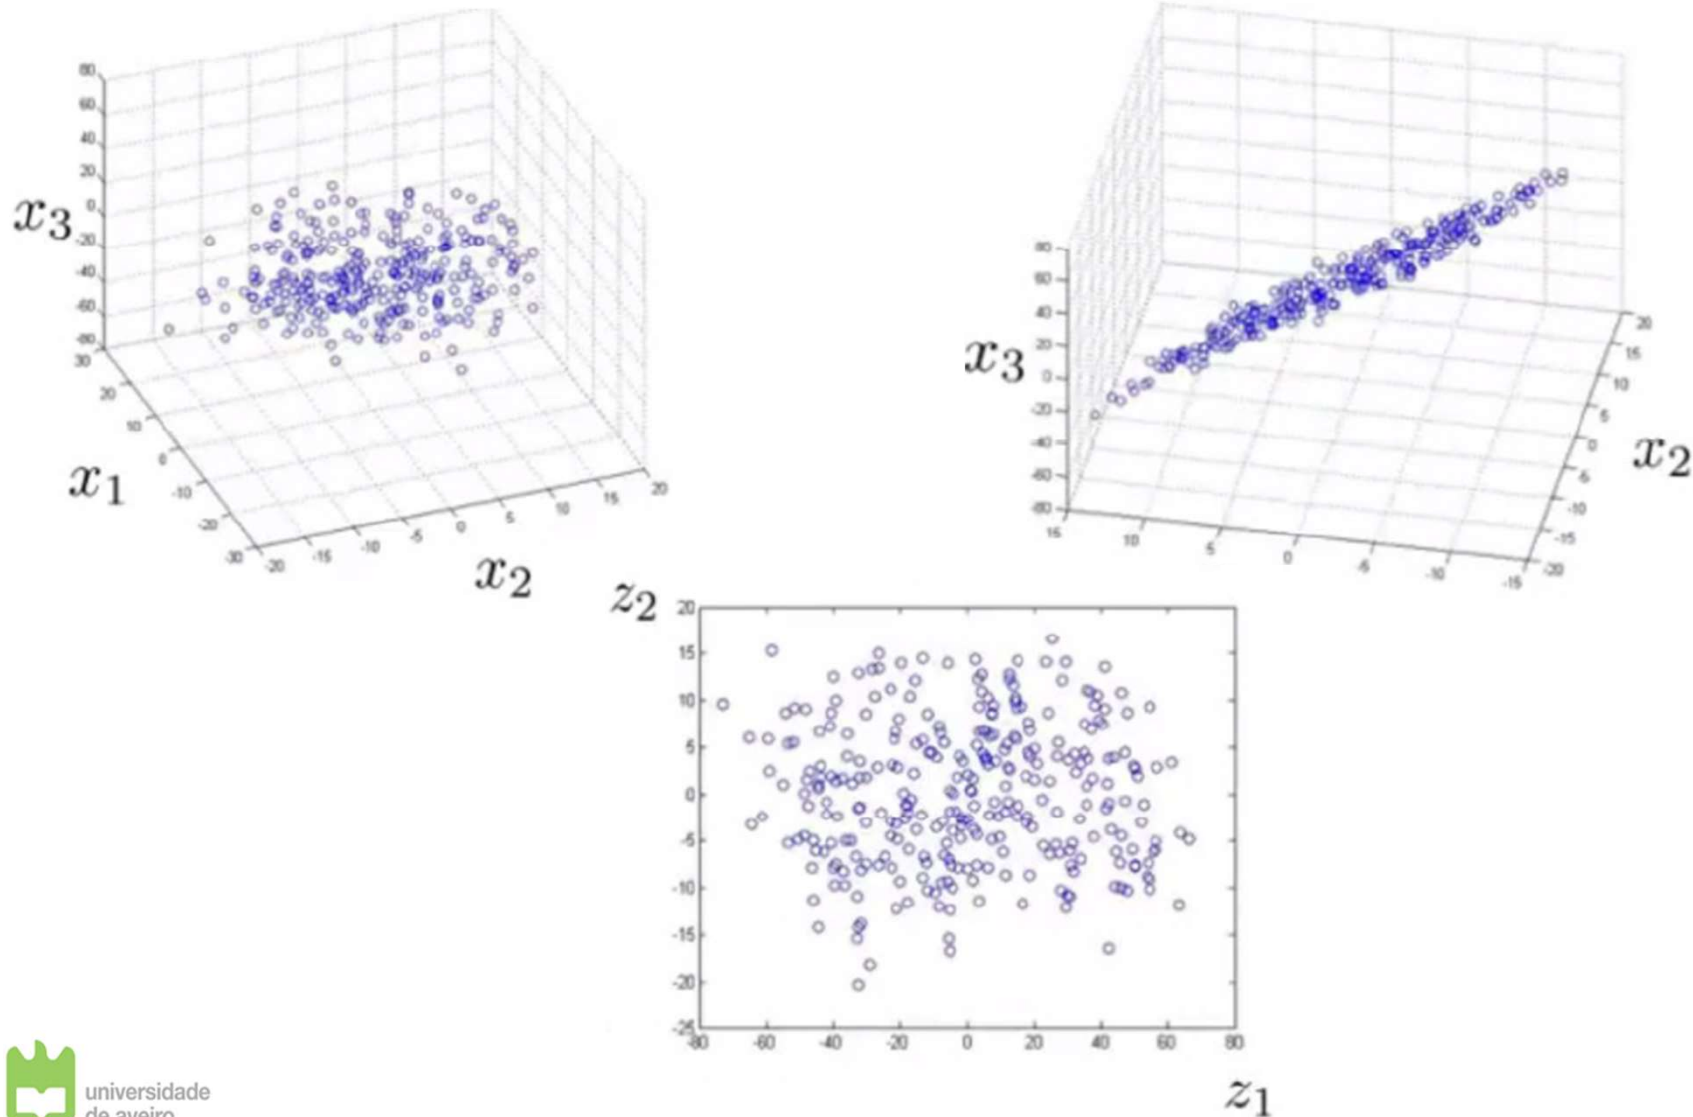




# DATA COMPRESSION

## Example: reduce data from 3D to 2D

If 3D data is located along a plane, the 3<sup>rd</sup> dimension is redundant



# DATA VISUALIZATION

**Example: Reduce data from high dimension to 2D or 3D**

Country	GDP (trillions of US\$)	Per capita GDP (thousands of intl. \$)	Human Develop- ment Index	Life expectancy	Poverty Index (Gini as percentage)	Mean household income (thousands of US\$)	...
Canada	1.577	39.17	0.908	80.7	32.6	67.293	...
China	5.878	7.54	0.687	73	46.9	10.22	...
India	1.632	3.41	0.547	64.7	36.8	0.735	...
Russia	1.48	19.84	0.755	65.5	39.9	0.72	...
Singapore	0.223	56.69	0.866	80	42.5	67.1	...
USA	14.527	46.86	0.91	78.3	40.8	84.3	...
...	...	...	...	...	...	...	...



# DATA VISUALIZATION

## Reduce data from high dimension to 2D or 3D

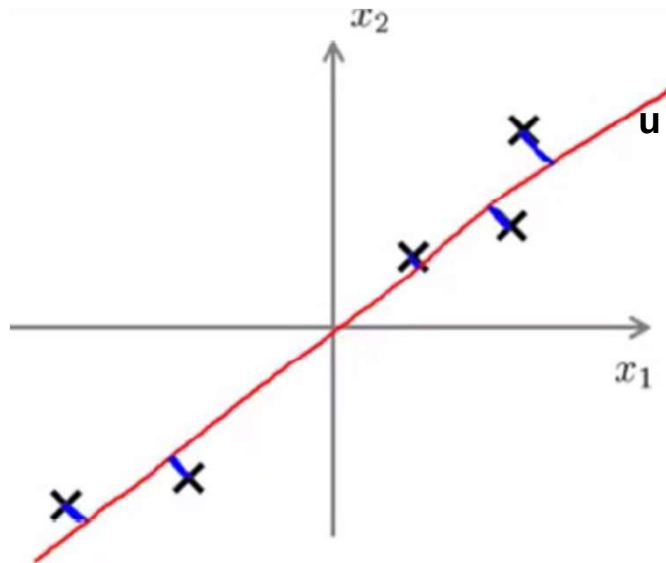
for example:  $z_1$ (country size/GDP) and  $z_2$  (GDP/per person)

Country	$z_1$	$z_2$
Canada	1.6	1.2
China	1.7	0.3
India	1.6	0.2
Russia	1.4	0.5
Singapore	0.5	1.7
USA	2	1.5
...	...	...

# PRINCIPAL COMPONENT ANALYSIS (PCA)

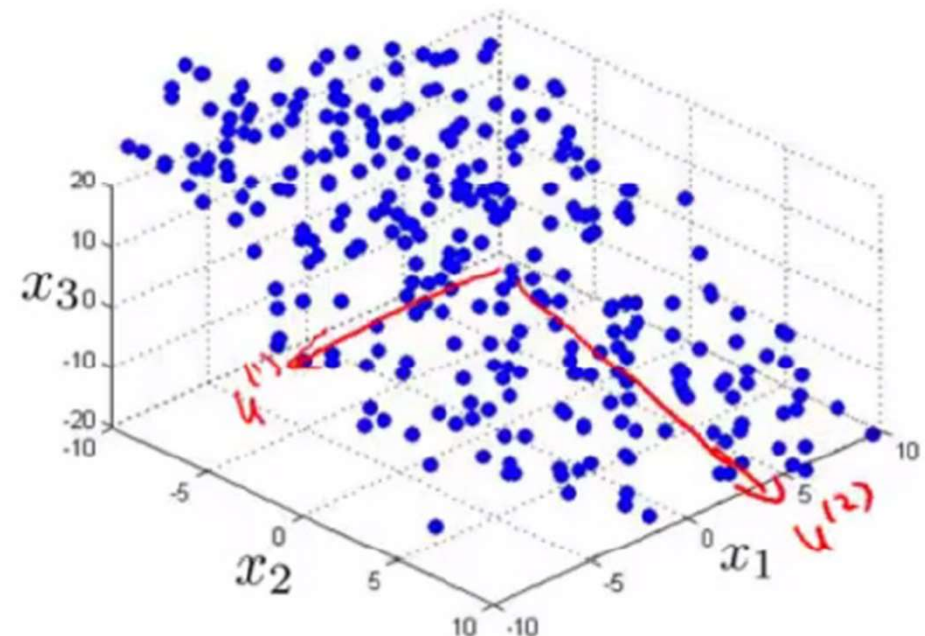
## Reduce from 2D to 1D:

find the best direction (vector  $u$ ) onto which to project data such that to minimize the projection error



## Reduce from 3D to 2D:

find the orientation of the best plane (vectors  $u_1, u_2$ ) onto which to project data such that to minimize the projection error

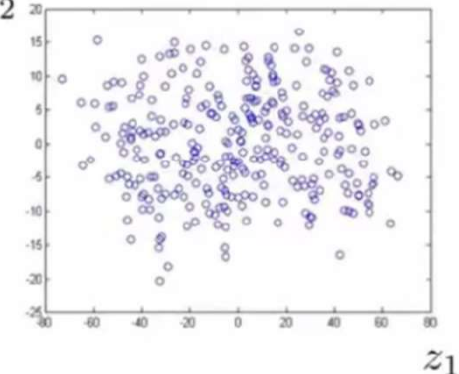
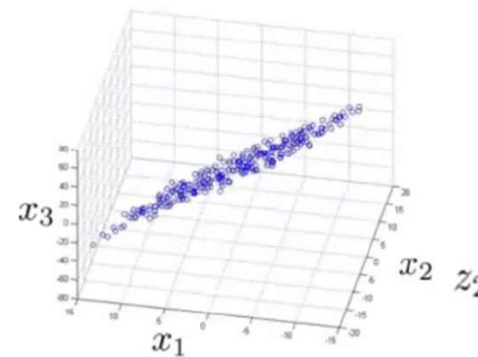
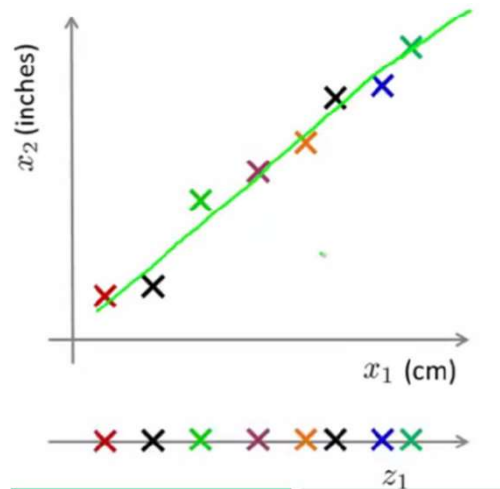


# PRINCIPAL COMPONENT ANALYSIS (PCA)

**Reduce from n-dimension to k-dimension ( $k < n$ ):**

**Two tasks:**

- Compute the coordinate system of the best k-dimensional surface (represented by vectors  $u_1, \dots, u_k$ ) onto which to project data such that to minimize the projection error.
- Compute the values of the projected data in the lower dimensional space (z values)



# PRINCIPAL COMPONENT ANALYSIS (PCA)

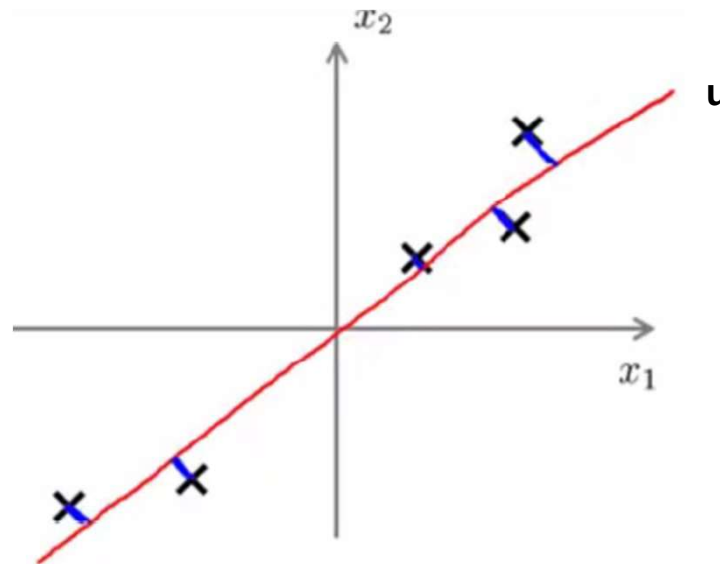
## PCA is not Linear Regression !!!

Linear Regression minimizes the vertical distance.

LR is a supervised learning method, it has a given output  $y$ .

PCA minimizes the orthogonal distance between the points and the projection surface (line, plane, etc.)

PCA is an unsupervised method, it has no output  $y$ .



# PCA – DATA PREPROCESSING (step 1)

Training set:  $x^{(1)}, x^{(2)}, \dots, x^{(m)}$

Preprocessing (feature scaling/mean normalization):

$$\mu_j = \frac{1}{m} \sum_{i=1}^m x_j^{(i)}$$

Replace each  $x_j^{(i)}$  with  $x_j - \mu_j$ .

**Thus, all features have zero mean !**

If the features have significantly different range of values, normalize them., e.g. in the interval  $[0,1]$  or  $[-1,1]$ .

## PCA – Singular Value Decomposition (step 2)

- Compute Covariance matrix of the mean normalized data matrix  $X$  (dimension  $m \times n$  -  $m$  examples,  $n$  features):

$$\mathbf{Cov} = \mathbf{X}^T * \mathbf{X}$$

- Compute Singular Value Decomposition(SVD) of Covariance matrix:

$$\mathbf{Cov} = \mathbf{U} * \mathbf{S} * \mathbf{V}$$

$U (n \times n)$  - matrix of eigenvectors:

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(n)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times n}$$

$S (n \times n)$  – diagonal matrix of singular values in decreasing order:

$$S_{n \times n} = \begin{bmatrix} S_{11} & 0 & \dots & 0 \\ 0 & S_{22} & \vdots & \vdots \\ \vdots & \vdots & \ddots & 0 \\ 0 & \dots & 0 & S_{nn} \end{bmatrix}$$

SVD is equivalent to eigen-values/eigen-vector decomposition.

Use linear algebra libraries to compute SVD =>  $[U, S, V] = \text{svd}(\text{Cov})$

# PCA

The projection vectors are the first  $k$  columns of  $U$  ( $k < n$ ):

$$U = \begin{bmatrix} | & | & & | \\ u^{(1)} & u^{(2)} & \dots & u^{(n)} \\ | & | & & | \end{bmatrix} \in \mathbb{R}^{n \times n} \quad Ureduce_{n \times k} = U(:, 1:k)$$

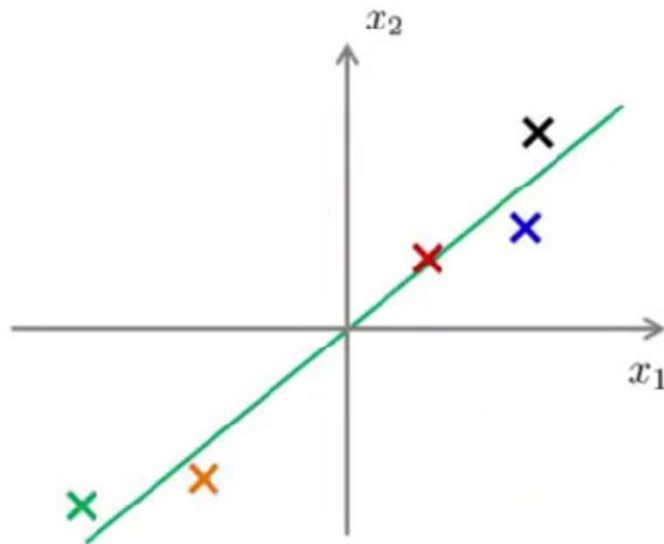
**Step 3:** Compute the new (projected) data matrix  $Z$  ( $m$  examples,  $k$  features):

$$Z_{m \times k} = X_{m \times n} * Ureduce_{n \times k}$$

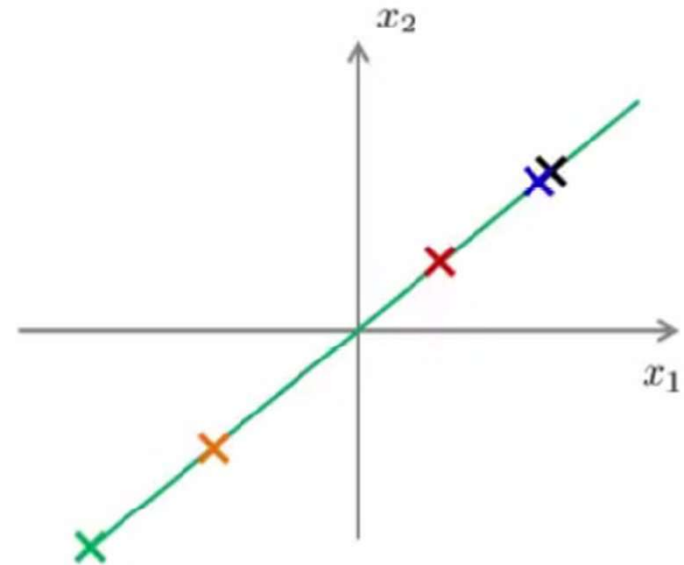
**Step 4:** Reconstruct data matrix  $X$  from the projected  $Z$  matrix :

$$X_{approx(m \times n)} = Z_{m \times k} * Ureduce_{k \times n}^T$$

# Reconstruction from compressed representation



$$z = U_{reduce}^T x$$



$$x_{approx} = U_{reduce} z$$



# Choosing $k$ (number of principal components)

**Average squared approximation error:**

$$\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{approx}^{(i)}\|^2$$

**Total data variation:**

$$\frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2$$

Typically, choose  $k$  to be smallest value so that

$$\frac{\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{approx}^{(i)}\|^2}{\frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2} \leq 0.01 \quad (1\%)$$

“99% of variance is retained”

(typically the desired retained variance is between 90-99%)

# Choosing $k$ (number of principal components)

**Algorithm 1** (highly inefficient, need to compute SVD several times):

Try PCA with  $k = 1$

Compute  $U_{reduce}, z^{(1)}, z^{(2)}, \dots, z^{(m)}, x_{approx}^{(1)}, \dots, x_{approx}^{(m)}$

Check if

$$\frac{\frac{1}{m} \sum_{i=1}^m \|x^{(i)} - x_{approx}^{(i)}\|^2}{\frac{1}{m} \sum_{i=1}^m \|x^{(i)}\|^2} \leq 0.01?$$

If not ok, try with  $k=2, 3, \dots$  until get desired retained variance

# Choosing $k$ (number of principal components)

**Algorithm 2 (Compute SVD once):**  $[U, S, V] = \text{svd}(\text{Cov})$

Increase  $k$  (starting from  $k=1$ ) until get the desired retained variance (e.g. 99 %):

$$\frac{\sum_{i=1}^k S_{ii}}{\sum_{i=1}^n S_{ii}} \geq 0.99$$

$$S_{n \times n} = \begin{bmatrix} S_{11} & 0 & \dots & 0 \\ 0 & S_{22} & \vdots & \vdots \\ \vdots & \vdots & \ddots & 0 \\ 0 & \dots & 0 & S_{nn} \end{bmatrix}$$

# PCA application

**To speed-up the supervised learning algorithm:**

- Get training set  $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$
- Run PCA to reduce  $x^{(i)}$  in dimension to get  $z^{(i)}$

**Run learning algorithm on the new training set with less features:**

$$\{(z^{(1)}, y^{(1)}), \dots, (z^{(m)}, y^{(m)})\}$$

**Note:** The projection matrix  $U_{reduce}$  is computed only once with the training data matrix  $X$ . This mapping is then applied to the cross validation and test examples.

**Bad use of PCA:** if the reason for applying PCA is not to speed-up learning but to prevent overfitting (less features less likely to overfit), better use regularization !!!